

Week 4 Exercises

Christopher Rodgers

April 6, 2024

Please complete all exercises below. You may use any library that we have covered in class. The data we will be using comes from the `tidyr` package, so you must use that.

- 1) Examine the `who` and `population` data sets that come with the `tidyr` library. the `who` data is not tidy, you will need to reshape the `new_sp_m014` to `newrel_f65` columns to long format retaining country, `iso2`, `iso3`, and year. The data in the columns you are reshaping contains patterns described in the details section below. You will need to assign three columns: diagnosis, gender, and age to the patterns described in the details.

Your tidy data should look like the following:

country	iso2	iso3	year	diagnosis	gender	age	count
1 Afghanistan	AF	AFG	1980	sp m	014	NA	2
2 Afghanistan	AF	AFG	1980	sp m	1524	NA	3
3 Afghanistan	AF	AFG	1980	sp m	2534	NA	4
4 Afghanistan	AF	AFG	1980	sp m	3544	NA	5
5 Afghanistan	AF	AFG	1980	sp m	4554	NA	6
6 Afghanistan	AF	AFG	1980	sp m	5564	NA	

Details The data uses the original codes given by the World Health Organization. The column names for columns five through 60 are made by combining `new_` to a code for method of diagnosis (`rel` = relapse, `sn` = negative pulmonary smear, `sp` = positive pulmonary smear, `ep` = extrapulmonary) to a code for gender (`f` = female, `m` = male) to a code for age group (`014` = 0-14 yrs of age, `1524` = 15-24 years of age, `2534` = 25 to 34 years of age, `3544` = 35 to 44 years of age, `4554` = 45 to 54 years of age, `5564` = 55 to 64 years of age, `65` = 65 years of age or older).

Note: use `data(who)` and `data(population)` to load the data into your environment. Use the arguments `cols`, `names_to`, `names_pattern`, and `values_to`. Your regex should be = (`"new_(.)_(.)_(.)"`)

<https://tidyr.tidyverse.org/reference/who.html>

```
data("who")
```

```
## Warning in data("who"): data set 'who' not found
```

```
data("population")
```

```
## Warning in data("population"): data set 'population' not found
```

```
library(dplyr)
```

```
## Warning: package 'dplyr' was built under R version 4.3.3
```

```
##
```

```
## Attaching package: 'dplyr'
```

```
## The following objects are masked from 'package:stats':
##
##   filter, lag
```

```
## The following objects are masked from 'package:base':
##
##   intersect, setdiff, setequal, union
```

```
library(tidyr)
```

```
library(ggplot2)
```

```
who_long <- who %>%
  pivot_longer(cols = 5:60,
               names_to = c('diagnosis', 'gender', 'age'),
               names_pattern = (regex= ("new_(.*)_(.)(.*)")),
               values_to = "count")
head(who_long)
```

```
## # A tibble: 6 x 8
##   country    iso2 iso3  year diagnosis gender age  count
##   <chr>      <chr> <chr> <dbl> <chr>      <chr> <chr> <dbl>
## 1 Afghanistan AF    AFG   1980 sp        m    014    NA
## 2 Afghanistan AF    AFG   1980 sp        m   1524    NA
## 3 Afghanistan AF    AFG   1980 sp        m   2534    NA
## 4 Afghanistan AF    AFG   1980 sp        m   3544    NA
## 5 Afghanistan AF    AFG   1980 sp        m   4554    NA
## 6 Afghanistan AF    AFG   1980 sp        m   5564    NA
```

- 2) There are two common keys between the data sets, with who as the left table, join the population data by country and year so that the population is available within the who dataset.

```
left_who <- who_long %>%
  left_join(population, by = c("country", "year"))
head(left_who)
```

```
## # A tibble: 6 x 9
##   country    iso2 iso3  year diagnosis gender age  count population
##   <chr>      <chr> <chr> <dbl> <chr>      <chr> <chr> <dbl>      <dbl>
## 1 Afghanistan AF    AFG   1980 sp        m    014    NA        NA
## 2 Afghanistan AF    AFG   1980 sp        m   1524    NA        NA
## 3 Afghanistan AF    AFG   1980 sp        m   2534    NA        NA
## 4 Afghanistan AF    AFG   1980 sp        m   3544    NA        NA
## 5 Afghanistan AF    AFG   1980 sp        m   4554    NA        NA
## 6 Afghanistan AF    AFG   1980 sp        m   5564    NA        NA
```

- 3) Split the age column into two columns, min age and max age. Notice that there is no character separator. Check the documentation with `?separate` to understand other ways to separate the age column. Keep in mind that 0 to 14 is coded as 014 (3 characters) and the other age groups are coded with 4 characters. 65 only has two characters, but we will ignore that until the next problem.

```
left_who <- left_who %>%
  separate('age', c('min_age', 'max_age'), sep = -2)

head(left_who)
```

```
## # A tibble: 6 x 10
##   country iso2 iso3 year diagnosis gender min_age max_age count population
##   <chr>    <chr> <chr> <dbl> <chr>    <chr> <chr> <chr> <dbl>    <dbl>
## 1 Afghanist~ AF   AFG   1980 sp      m      0     14      NA      NA
## 2 Afghanist~ AF   AFG   1980 sp      m     15     24      NA      NA
## 3 Afghanist~ AF   AFG   1980 sp      m     25     34      NA      NA
## 4 Afghanist~ AF   AFG   1980 sp      m     35     44      NA      NA
## 5 Afghanist~ AF   AFG   1980 sp      m     45     54      NA      NA
## 6 Afghanist~ AF   AFG   1980 sp      m     55     64      NA      NA
```

- 4) Since we ignored the 65+ group in the previous problem we will fix it here. If you examine the data you will notice that 65 was placed into the max_age column and there is no value for min_age for those records. To fix this use mutate() in order to replace the blank value in the min_age column with the value from the max_age column and another mutate to replace the 65 in the max column with an Inf. Be sure to keep the variables as character vectors.

```
left_who <- left_who %>%
  mutate(min_age = replace(min_age, min_age == '', '64')) %>%
  mutate(max_age = replace(max_age, min_age == "65", "inf"))

head(left_who)
```

```
## # A tibble: 6 x 10
##   country iso2 iso3 year diagnosis gender min_age max_age count population
##   <chr>    <chr> <chr> <dbl> <chr>    <chr> <chr> <chr> <dbl>    <dbl>
## 1 Afghanist~ AF   AFG   1980 sp      m      0     14      NA      NA
## 2 Afghanist~ AF   AFG   1980 sp      m     15     24      NA      NA
## 3 Afghanist~ AF   AFG   1980 sp      m     25     34      NA      NA
## 4 Afghanist~ AF   AFG   1980 sp      m     35     44      NA      NA
## 5 Afghanist~ AF   AFG   1980 sp      m     45     54      NA      NA
## 6 Afghanist~ AF   AFG   1980 sp      m     55     64      NA      NA
```

- 5) Find the count per diagnosis for males and females.

See ?sum for a hint on resolving NA values.

```
count_diagnosis <- left_who %>%
  group_by(diagnosis, gender) %>%
  summarise(count = sum(!is.na(count)), .groups = "drop")

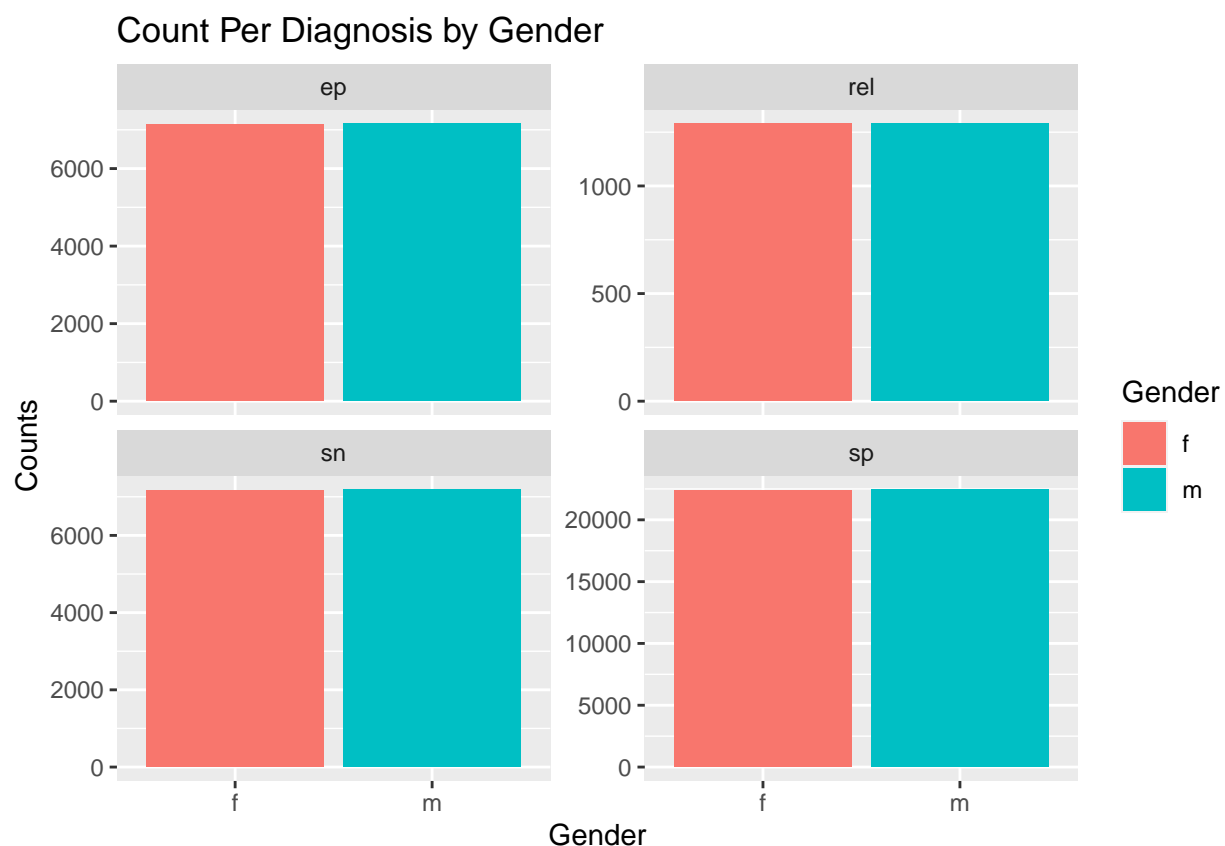
head(count_diagnosis)
```

```
## # A tibble: 6 x 3
##   diagnosis gender count
##   <chr>    <chr> <int>
## 1 ep      f      7143
```

```
## 2 ep      m      7161
## 3 rel     f      1290
## 4 rel     m      1290
## 5 sn      f      7152
## 6 sn      m      7190
```

- 6) Now create a plot using ggplot and geom_col where your x axis is gender, your y axis represents the counts, and facet by diagnosis. Be sure to give your plot a title and resolve the axis labels.

```
ggplot(count_diagnosis, aes (x = gender, y = count, fill = gender)) +
  geom_col(position = "dodge") +
  facet_wrap(~ diagnosis, scales = "free_y", ncol = 2) +
  labs(x = "Gender", y = "Counts", fill = "Gender", title = "Count Per Diagnosis by Gender")
```



- 7) Find the percentage of population by year, gender, and diagnosis. Be sure to remove rows containing NA values.

```
population_who <- left_who %>%
  drop_na() %>%
  group_by(year, gender, diagnosis) %>%
  summarise(percentage = (100 * (sum(count)) / (sum(population))))
```

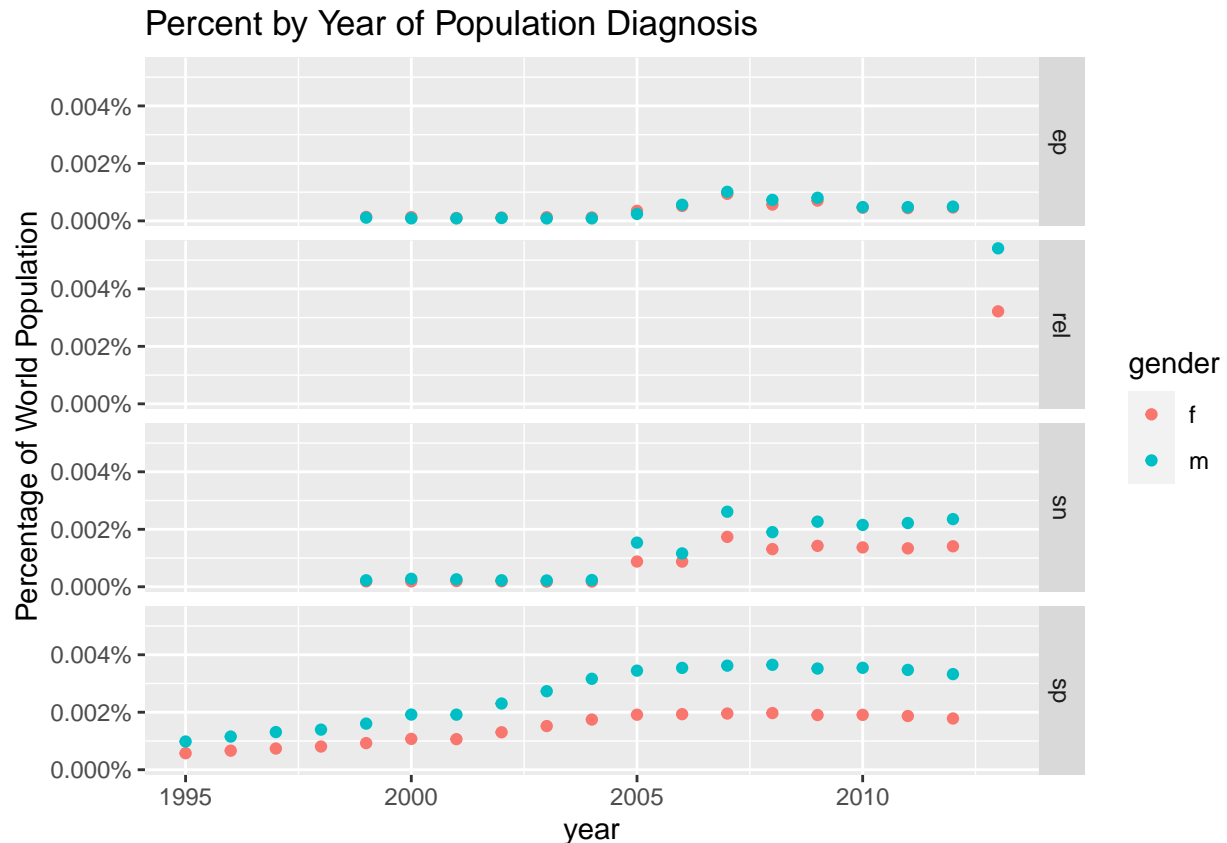
```
## 'summarise()' has grouped output by 'year', 'gender'. You can override using
## the '.groups' argument.
```

```
print(population_who)
```

```
## # A tibble: 94 x 4
## # Groups:   year, gender [38]
##   year gender diagnosis percentage
##   <dbl> <chr> <chr>         <dbl>
## 1 1995 f      sp           0.000574
## 2 1995 m      sp           0.000982
## 3 1996 f      sp           0.000663
## 4 1996 m      sp           0.00115
## 5 1997 f      sp           0.000737
## 6 1997 m      sp           0.00131
## 7 1998 f      sp           0.000807
## 8 1998 m      sp           0.00139
## 9 1999 f      ep           0.000138
## 10 1999 f      sn           0.000188
## # i 84 more rows
```

- 8) Create a line plot in ggplot where your x axis contains the year and y axis contains the percent of world population. Facet this plot by diagnosis with each plot stacked vertically. You should have a line for each gender within each facet. Be sure to format your y axis and give your plot a title.

```
ggplot(population_who) +
  geom_point(aes (x = year, y = (percentage/100), color=gender)) +
  facet_grid(diagnosis ~ .) +
  scale_y_continuous(labels = scales :: percent_format()) +
  labs(title = 'Percent by Year of Population Diagnosis', y= 'Percentage of World Population')
```



- 9) Now unite the min and max age variables into a new variable named `age_range`. Use a '-' as the separator.

```
left_who <- left_who %>%
  unite(age_range, min_age, max_age, sep = '-')
head(left_who)
```

```
## # A tibble: 6 x 9
##   country iso2 iso3 year diagnosis gender age_range count population
##   <chr>    <chr> <chr> <dbl> <chr>    <chr>    <chr>    <dbl>    <dbl>
## 1 Afghanistan AF AFG 1980 sp      m      0-14      NA      NA
## 2 Afghanistan AF AFG 1980 sp      m     15-24      NA      NA
## 3 Afghanistan AF AFG 1980 sp      m     25-34      NA      NA
## 4 Afghanistan AF AFG 1980 sp      m     35-44      NA      NA
## 5 Afghanistan AF AFG 1980 sp      m     45-54      NA      NA
## 6 Afghanistan AF AFG 1980 sp      m     55-64      NA      NA
```

- 10) Find the percentage contribution of each age group by diagnosis. You will first need to find the count of all diagnoses then find the count of all diagnoses by age group. Join the former to the later and calculate the percent of each age group. Plot these as a `geom_col` where the x axis is the diagnosis, y axis is the percent of total, and faceted by age group.

```
count_all <- left_who %>%
  group_by(diagnosis) %>%
  summarise(total_count = sum(count, na.rm = TRUE))

head(count_all)
```

```
## # A tibble: 4 x 2
##   diagnosis total_count
##   <chr>         <dbl>
## 1 ep           1986179
## 2 rel          3220572
## 3 sn           6279527
## 4 sp           31911240
```

```
by_age_group <- left_who %>%
  count(diagnosis, age_range, name = "age_range_count")

head(by_age_group)
```

```
## # A tibble: 6 x 3
##   diagnosis age_range age_range_count
##   <chr>      <chr>         <int>
## 1 ep       0-14             14480
## 2 ep       15-24             14480
## 3 ep       25-34             14480
## 4 ep       35-44             14480
## 5 ep       45-54             14480
## 6 ep       55-64             14480
```

```
by_diagnosis_age <- count_all %>%
  left_join(by_age_group, by= 'diagnosis')

head(by_diagnosis_age)
```

```
## # A tibble: 6 x 4
##   diagnosis total_count age_range age_range_count
##   <chr>         <dbl> <chr>         <int>
## 1 ep           1986179 0-14             14480
## 2 ep           1986179 15-24             14480
## 3 ep           1986179 25-34             14480
## 4 ep           1986179 35-44             14480
## 5 ep           1986179 45-54             14480
## 6 ep           1986179 55-64             14480
```

```
by_diagnosis_age <- by_diagnosis_age %>%
  mutate(percentage = age_range_count / total_count *100)

head(by_diagnosis_age)
```

```
## # A tibble: 6 x 5
##   diagnosis total_count age_range age_range_count percentage
```

##	<chr>	<dbl>	<chr>	<int>	<dbl>
## 1	ep	1986179	0-14	14480	0.729
## 2	ep	1986179	15-24	14480	0.729
## 3	ep	1986179	25-34	14480	0.729
## 4	ep	1986179	35-44	14480	0.729
## 5	ep	1986179	45-54	14480	0.729
## 6	ep	1986179	55-64	14480	0.729

```
ggplot(by_diagnosis_age, aes(x = diagnosis, y = percentage, fill = age_range)) +
  geom_col(position = "dodge") +
  facet_wrap(~age_range, scales = "free_y", ncol = 2) +
  labs(x = "Diagnosis", y = "Percent of Total", fill = "Age_Range", title = "Age Group Percentage Total")
```

